

Single Node Optimisation

Optimising with the Compiler



Reusing this material



This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License.

<https://creativecommons.org/licenses/by-nc-sa/4.0/>

This means you are free to copy and redistribute the material and adapt and build on the material under the following terms: You must give appropriate credit, provide a link to the license and indicate if changes were made. If you adapt or build on the material you must distribute your work under the same license as the original.

Note that this presentation contains images owned by others. Please seek their permission before reusing these images.

- Unless we write assembly code, we are always using a compiler.
- Modern compilers are (quite) good at optimisation
 - memory optimisations are an exception
- Usually much better to get the compiler to do the optimisation.
 - avoids machine-specific coding
 - compilers break codes much less often than humans
- Even modifying code can be thought of as “helping the compiler”.

Compiler flags

- Typical compiler has hundreds of flags/options.
 - most are never used
 - many are not related to optimisation
- Most compilers have flags for different levels of general optimisation.
 - **-O1, -O2, -O3,....**
- When first porting code, switch optimisation off.
 - only when you are satisfied that the code works, turn optimisation on, and test again.
 - but don't forget to use them!
 - also don't forget to turn off debugging, bounds checking and profiling flags...

Compiler flags (cont.)

- Note that highest levels of optimisation may
 - break your code.
 - give different answers, by bending standards.
 - make your code go slower.
- Always read documentation carefully.
- Isolate routines and flags which cause the problem.
 - binary chop
 - one routine per file may help

Compiler flags (cont.)



- Many compilers are designed for an instruction set architecture, not one machine.
 - flags to target ISA versions, processor versions, cache configurations
 - defaults may not be optimal, especially if cross-compiling
- Some optimisation flags may not be part of **-On**
 - check documentation
 - use sparingly (may only be beneficial in some cases)

- A mechanism for giving additional information to the compiler, e.g.
 - values of variables (e.g. loop trip counts)
 - independence of loop iterations
 - independence of index array elements
 - aliasing properties
- Appear as comments (Fortran), or pre-processor pragmas (C)
 - don't affect portability

Incremental compilation



- Compilers can only work with the limited information available to them.
- Most compilers compile code in an incremental fashion
 - Each source file is compiled independently of each other.
 - Most compilers ignore all source files other than those specified on the command line (or implicitly referenced via search paths, e.g. include files)
 - Routines from other source files treated as “black-boxes”
 - Make worst case assumptions based on routine prototype.
- You can help by providing more information
 - Information in routine prototypes
 - INTENT, PURE, const, etc.
 - Compiler hints
 - Command line flags

- When flags and hints don't solve the problem, we will have to resort to code modification.
- Be aware that this may
 - introduce bugs.
 - make the code harder to read/maintain.
 - only be effective on certain architectures and compiler versions.
- Try to think about
 - what optimisation the compiler is failing to do
 - what additional information can be provided to compiler
 - how can rewriting help

- How can we work out what the compiler has done?
 - eyeball assembly code
 - use diagnostics flags
- Increasingly difficult to work out what actually occurred in the processor.
 - superscalar, out-of-order, speculative execution
- Can estimate expected performance
 - count flops, load/stores, estimate cache misses
 - compare actual performance with expectations

Locals and globals

- Compiler analysis is more effective with local variables
- Has to make worst case assumptions about global variables
- Globals could be modified by any called procedure (or by another thread).
- Use local variables where possible
- Automatic variables are stack allocated: allocation is essentially free.
- In C, use file scope globals in preference to externals

Conditionals

- Even with sophisticated branch prediction hardware, branches are bad for performance.
- Try to avoid branches in innermost loops.
 - if you can't eliminate them, at least try to get them out of the critical loops.

```
do i=1,k
  if (n .eq. 0) then
    a(i) = b(i) + c
  else
    a(i) = 0.
  endif
end do
```



```
if (n .eq. 0) then
  do i=1,k
    a(i) = b(i) + c
  end do
else
  do i=1,k
    a(i) = 0.
  end do
endif
```

- A little harder for the compiler.....

```
do i=1,k
  if (i .le. j) then
    a(i) = b(i) + c
  else
    a(i) = 0.
  endif
end do
```



```
do i=1,j
  a(i) = b(i) + c
end do
do i = j+1,k
  a(i) = 0.
end do
```

- Performance can be affected by choice of data types
 - often a difference between 32-bit and 64-bit arithmetic (integer and floating point).
 - complicated by trade-offs with memory usage and cache hit rates
- Avoid unnecessary type conversions
 - e.g. int to long, float to double
 - N.B. some type conversions are implicit
 - However sometimes better than the alternative e.g.
 - Use DP reduction variable rather than increase array precision.

- Compilers are generally good at Common Subexpression Elimination.
- A couple of cases where they might have trouble:

Different order of operands

```
d = a + c  
e = a + b + c
```

Function calls

```
d = a + func(c)  
e = b + func(c)
```

CSE including function calls.

- To extract a CSE containing a function call the compiler has to be sure of various things:
 - The function always returns the same value for the same input.
 - The function does not cause any side effects that would be affected by changing the number of times the function is called:
 - Modifying its inputs.
 - Changing global data.
- Need to be very careful with function prototypes to allow compiler to know this.

- Most compilers make a reasonable job of register allocation.
 - But only limited number available.
- Can have problems in some cases:
 - loops with large numbers of temporary variables
 - such loops may be produced by inlining or unrolling
 - array elements with complex index expressions
 - can help compiler by introducing explicit scalar temporaries, most compilers will use a register for an explicit scalar in preference to an implicit CSE.

```
for (i=0;i<n;i++){  
    b[i] += a[c[i]];  
    c[i+1] = 2*i;  
}
```



```
tmp = c[0];  
for (i=0;i<n;i++){  
    b[i] += a[tmp];  
    tmp = 2*i;  
    c[i+1] = tmp;  
}
```

- If compiler runs out of registers it will generate spill code.
 - store a value and then reload it later on
- Examine your source code and count how many loads/stores are required
- Compare with assembly code
- May need to distribute loops

- Loop unrolling and software pipelining are two of the most important optimisations for scientific codes on modern RISC processors.
- Compilers generally good at this.
- If compiler fails, usually better to try and remove the impediment, rather than unroll by hand.
 - cleaner, more portable, better performance
- Compiler has to determine independence of iterations

- Loops with small bodies generate small basic blocks of assembly code
 - lot of dependencies between instructions
 - high branch frequency
 - little scope for good instruction scheduling
- Loop unrolling is a technique for increasing the size of the loop body
 - gives more scope for better schedules
 - reduces branch frequency
 - make more independent instructions available for multiple issue.

Loop unrolling

- Replace loop body by multiple copies of the body
- Modify loop control
 - take care of arbitrary loop bounds
- Number of copies is called **unroll factor**

Example:

```
do i=1,n
  a(i)=b(i)+d*c(i)
end do
```



```
do i=1,n-3,4
  a(i)=b(i)+d*c(i)
  a(i+1)=b(i+1)+d*c(i+1)
  a(i+2)=b(i+2)+d*c(i+2)
  a(i+3)=b(i+3)+d*c(i+3)
end do
do j = i,n
  a(j)=b(j)+d*c(j)
end do
```

- Remember that this is in fact done by the compiler at the IR or assembly code level.
- If the loop iterations are independent, then we end up with a larger basic block with relatively few dependencies, and more scope for scheduling.
 - also reduce no. of compare and branch instructions
- Choice of unroll factor is important (usually 2,4,8)
 - if factor is too large, can run out of registers
- Cannot unroll loops with complex flow control
 - hard to generate code to jump out of the unrolled version at the right place

Outer loop unrolling

- If we have a loop nest, then it is possible to unroll one of the outer loops instead of the innermost one.
- Can improve locality.

```
do i=1,n
  do j=1,m
    a(i,j)=c*d(j)
  end do
end do
```

2 loads for 1 flop



```
do i=1,n,4
  do j=1,m
    a(i,j)=c*d(j)
    a(i+1,j)=c*d(j)
    a(i+2,j)=c*d(j)
    a(i+3,j)=c*d(j)
  end do
end do
```

5 loads for 4 flops

- Variable expansion can help break dependencies in unrolled loops
 - improves scheduling opportunities
- Close connection to reduction variables in parallel loops

```
for (i=0,i<n,i++){  
    b+=a[i];  
}
```

unroll

```
for (i=0,i<n,i+=2){  
    b+=a[i];  
    b+=a[i+1];  
}
```

expand b

```
for (i=0,i<n,i+=2){  
    b1+=a[i];  
    b2+=a[i+1];  
}  
b=b1+b2;
```

Register renaming

- Registers may be reused within a basic block introducing unnecessary dependencies.
- Using two (or more) different registers can preserve program correctness, but allow more scheduling flexibility
 - Some CPUs perform register rename and reschedule in hardware, this can utilise additional registers not visible to compiler.

```
add %f2,1,%f1  
st [%o1],f1  
add %f3,2,%f1  
st [%o2],f1
```

rename

```
add %f2,1,%f1  
st [%o1],f1  
add %f3,2,%f27  
st [%o2],f27
```

reschedule

```
add %f2,1,%f1  
add %f3,2,%f27  
st [%o1],f1  
st [%o2],f27
```

- Problem with scheduling small loop bodies is that there are dependencies between instructions in the basic block.
- Potentially possible to start executing instructions from the next iteration before current one is finished.
- Idea of software pipelining is to construct a basic block that contains instructions from **different** loop iterations.
 - fewer dependencies between instructions in block
 - needs additional code at start and end of loop

Software pipelining

```
for (i=0;i<n;i++){  
    a(i) += b;  
}
```



```
for (i=0;i<n;i++){  
    t1 = a(i);    //L i  
    t2 = b + t1;  //A i  
    a(i) = t2;    //S i  
}
```



```
//prologue  
t1 = a(0);    //L 0  
t2 = b + t1;  //A 0  
t1 = a(1);    //L 1  
  
for (i=0;i<n-2;i++){  
    a(i) = t2;    //S i  
    t2 = b + t1;  //A i+1  
    t1 = a(i+2);  //L i+2  
}  
  
//epilogue  
a(n-2) = t2;    //S n-2  
t2 = b + t1;    //A n-1  
a(n-1) = t2;    //S n-1
```

Instruction level

```
L: ld    [%r1],%f0
    fadd  f0,f1,f2
    st    [%r1],f2
    add   %r1,4,%r1
    cmp   %r1,%r3
    bg    L
    nop
```



```
ld    [%r1],%f0
fadd   f0,f1,f2
ld     [%r1+4],%f0

L: st    [%r1],f2
    fadd  f0,f1,f2
    ld    [%r1+8],%f0
    cmp   %r1,%r3-8
    bg    L
    add   %r1,4,%r1

    st    [%r1],f2
    add   %r1,4,%r1
    fadd  f0,f1,f2
    st    [%r1],f2
```

st must wait for **fadd**
to complete: pipeline stall
for data hazard

Impediments to unrolling

- Function calls
 - except in presence of good interprocedural analysis and inlining
- Conditionals
 - especially control transfer out of the loop
 - lose most of the benefit anyway as they break up the basic block.
- Pointer/array aliasing
 - compiler can't be sure different values don't overlap in memory

Example

```
for (i=0;i<ip;i++){  
    a[indx[i]] += c[i] * a[ip];  
}
```

- Compiler doesn't know that `a[indx[i]]` and `a[ip]` don't overlap
- Could try hints
 - tell compiler that `indx` is a permutation
 - tell compiler that it is OK to unroll
- Or could rewrite:

```
tmp = a[ip];  
for (i=0;i<ip;i++){  
    a[indx[i]] += c[i] * tmp;  
}
```

Inlining

- Compilers very variable in their abilities
- Hand inlining possible
 - very ugly (slightly less so if done via pre-processor macros)
 - causes code replication
- Compiler has to know where the source of candidate routines is.
 - sometimes done by compiler flags
 - easier for routines in the same file
 - try compiling multiple files at the same time
- Very important for OO code
 - OO design encourages methods with very small bodies
 - inline keyword in C++ can be used as a hint

Multiple Optimisation steps

- Sometimes multiple optimisation steps are required.
 - Multiple levels of in-lining.
 - In-lining followed by loop un-rolling followed by CSE.
- The compiler may not be able to perform all steps at the same time
 - You may be able to help the compiler by performing some of the steps by hand.
 - Look for the least damaging code change that allows the compiler to complete the rest of the necessary changes.
 - Ideally try each step in isolation before attempting to combine hand-optimisations.

General Cray Fortran Compiler Flags



- Optimisation Options

- **-O2** optimal flags [enabled by default]
- **-O3** aggressive optimization
- **-O ipaN** inlining, $N=0-5$ [default $N=3$]

- Create listing files with optimization info

- **-hlist=a** creates a listing file with all optimization info
- **-hlist=m** produces a source listing with loopmark information

- Parallelization Options

- **-h omp** Recognize OpenMP directives
- **-h threadN** control the compilation and optimization of OpenMP directives, $N=0-3$ [default $N=2$]

➔ More info: `man crayftn`

Recommended CCE Fortran Compilation Options | epcc |

- Use default optimization level (**-O2**) as a starting point
 - It's the equivalent of most other compilers -O3 or -fast
- Use **-O3 -hfp3** if the application runs cleanly with these options
 - **-O3** only gives you slightly more than the default **-O2**
 - Cray also test this thoroughly
 - **-hfp3** gives you a lot more floating point optimization (default is **-hfp2**)
- If an application is intolerant of floating point reordering, try a lower **-hfp2** number
 - Try **-hfp1** first, only **-hfp0** if absolutely necessary (**-hfp4** is the maximum)
 - Might be needed for tests that require strict IEEE conformance
 - Or applications that have 'validated' results from a different compiler
- Do not use too aggressive optimizations , e.g. **-hfp4** without a very good reason
 - Higher numbers are not always correlated with better performance

General Cray C/C++ Compiler Flags



- Optimisation Options

- **-O3**

standard optimisation

- **-Ofast**

includes more aggressive FP optimisations

- **-flto**

link-time optimisations

- Create listing files with optimization info

- **-fsave-loopmark**

creates a listing file with all optimization info

- Parallelization Options

- **-fopenmp**

Recognize OpenMP directives

➔ More info: `man craycc`, `man crayCC`

Recommended CCE C/C++ Compilation Options



- Use **-O3** as a starting point for optimisation
- Use **-Ofast -flto** if the application runs cleanly with these options
 - **-ffa=N** gives you more control (N=0-4, **-Ofast** implies **-ffa=3**)
 - If an application is intolerant of floating point reordering, try a lower **-ffa** number
- Do not use too aggressive optimizations , e.g. **-ffa=4** without a very good reason
 - Higher numbers are not always correlated with better performance

- More or less all optimizations and features provided by CCE are available in AMD and GNU compilers
 - GNU compiler serves a wide range of users & needs
 - Default compiler with Linux, some people only test with GNU
 - **GNU defaults are conservative (e.g. -O1)**
 - **-O3** includes vectorization and most inlining
 - Performance users set additional options
 - AMD compiler is typically more aggressive in the optimizations
 - **AMD defaults are more aggressive (e.g. -O2)**, to give better performance “out-of-the-box”
 - Includes vectorization; some loop transformations such as unrolling; inlining within source file
 - Options to scale back optimizations for better floating-point reproducibility, easier debugging, etc.
 - Additional options for optimizations less sure to benefit all applications
 - **CCE is even more aggressive** in the optimizations by default
 - Better inlining and vectorization
 - Aggressive floating-point optimizations
- Bottom line: GNU users probably have to specify higher optimisation levels

Summary

- Remember compiler is always there.
- Try to help compiler, rather than do its job for it.
- Use flags and hints as much as possible
- Minimise code modifications