



Single Node Optimisation Vectorisation



Reusing this material



This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License.

<https://creativecommons.org/licenses/by-nc-sa/4.0/>

This means you are free to copy and redistribute the material and adapt and build on the material under the following terms: You must give appropriate credit, provide a link to the license and indicate if changes were made. If you adapt or build on the material you must distribute your work under the same license as the original.

Note that this presentation contains images owned by others. Please seek their permission before reusing these images.



Vector Instructions (Vectorisation)



- Modern CPUs can perform multiple operations each cycle
 - Use special SIMD (Single Instruction Multiple Data) instructions
 - e.g. SSE, AVX
 - Operate on a "vector" of data
 - typically 2 or 4 double precision floats (on AMD Rome)
 - But can be up to 8 per FPU
 - Potentially gives speedup in floating point operations
 - Usually only one loop is vectorisable in loop nest
 - And most compilers only consider inner loop




Vectorisation


- Same operation on multiple data items
 - Wide registers
 - SIMD needed to approach FLOP peak performance, but your code must be capable of vectorisation

- **x86 SIMD instruction sets:**


- SSE: register width = 128 Bit
 - 2 double precision floating point operands
- AVX: register width = 256 Bit
 - 4 double precision floating point operands



```
for(i=0;i<N;i++){
    a[i] = b[i] + c[i];
}
do i=1,N
    a(i) = b(i) + c(i)
end do
```



AVX2/AVX256



- Rome processor has two AVX256 vector units per core
 - Symmetrical units
 - Only one supports some of the legacy stuff (x87, MMX, some of the SSE stuff)
 - Vector instructions have a latency of 6 cycles

- Optimising compilers will use vector instructions
 - Relies on code being vectorisable
 - ...or in a form that the compiler can convert to be vectorisable
 - Some compilers are better at this than others
 - But there are some general guidelines about what is likely to work...

When does the compiler vectorize

- What can be vectorized
 - Only loops
- Usually only one loop is vectorisable in loopnest
 - And most compilers only consider inner loop
- Optimising compilers will use vector instructions
 - Relies on code being vectorisable
 - Or in a form that the compiler can convert to be vectorisable
 - Some compilers are better at this than others
- Check the compiler output listing and/or assembler listing
 - Look for packed AVX/AVX2/AVX512 instructions

i.e. Instructions using registers `zmm0–zmm31` (512-bit) `ymm0–ymm31` (256-bit) `xmm0–xmm31` (128-bit)

Instructions like `vaddps`, `vmulps`, etc...

Requirements for vectorisation



- Loops must have determinable (at run time) trip count
 - rules out most while loops
- Loops must not contain function/subroutine calls
 - unless the call can be inlined by the compiler
 - maths library functions usually OK (but not with GNU compilers)
- Loops must not contain branches or jumps
 - guarded assignments may be OK
 - e.g. `if (a[i] != 0.0) b[i] = c * a[i];`
- Loop trip counts needs to be long, or else a multiple of the vector length



- Loops must not have dependencies between iterations
 - reductions usually OK, e.g. `sum += a[i];`
 - avoid induction variables e.g. `indx += 3;`
 - use **restrict**
 - may need to tell the compiler if it can't work it out for itself
- Aligned data is best
 - e.g. AVX vector loads/stores operate most effectively on 32-bytes aligned address
 - need to either let the compiler align the data....
 - ..or tell it what the alignment is
- Unit stride through memory is best

Compilers



- Cray (C) and AMD compilers requires
 - Optimisation enabled (generally is by default)
 - `-O2`
 - To know what hardware it's compiling for
 - `-march=znver2`
 - This is added automatically for you on ARCHER2
 - Can disable vectorisation
 - `-fno-vectorize`
 - Useful for checking performance
 - Cray compiler will provide vectorisation information
 - `-Rpass-missed=loop-vectorize -Rpass-analysis=loop-vectorize`
 - Other compilers information
 - Cray Fortran: `-hlist=a`
 - GNU: `-fdump-tree-vect-all=<filename>`



Did my loop get vectorised?



- Always check the compiler output to see what it did
 - CCE: `-hlist=a`
 - GNU: `-fdump-tree-vect-all=<filename>`
 - AMD: `-Rpass-missed=loop-vectorize -Rpass-analysis=loop-vectorize`
 - or (for the hard core) check the assembler generated
 - Look to see which registers are in use.
- Clues from CrayPAT's HWPC measurements
 - `export PAT_RT_HWPC=13` or `14` # Floating point operations SP,DP
 - Complicated, but look for ratio of operations/instructions > 1
 - expect 4 for pure AVX with double precision floats



Did my loop get vectorised?



- GNU offers other options for checking:
- `-fopt-info`
- `-fopt-info-all`
- `-O3 -fopt-info-missed=missed.all`
- `-O2 -ftree-vectorize -fopt-info-vec-missed`
- `-fopt-info-loop-optimized`



Helping vectorisation

- Does the loop have dependencies?
 - information carried between iterations
 - e.g. `counter: total = total + a(i)`
 - No:
 - Tell the compiler that it is safe to vectorise
 - Yes:
 - Rewrite code to use algorithm without dependencies, e.g.
 - promote loop scalars to vectors (single dimension array)
 - use calculated values (based on loop index) rather than iterated counters, e.g.
 - Replace: `count = count + 2; a(count) = ...`
 - By: `a(2*i) = ...`
 - move `if` statements outside the inner loop
 - may need temporary vectors to do this (otherwise use masking operations)
- Is there a good reason for this?
 - There is an overhead in setting up vectorisation; maybe it's not worth it
 - Could you unroll inner (or outer) loop to provide more work?

Vectorisation example

- Compiler cannot easily vectorise:
 - Loops with pointers
 - Non-unit stride loops
 - Funny memory patterns
 - Unaligned data accesses
 - Conditionals/Function calls in loops
 - Data dependencies between loop iterations
 -

```
int *loop_size;
void problem_function(float *data1, float *data2, float *data3, int *index){
    int i,j;
    for(i=0;i<*loop_size;i++){
        j = index[i];
        data1[j] = data2[i] * data3[i];
    }
}
```

Vectorisation example



- Can help compiler
 - Tell it loops are independent
 - `#pragma clang loop vectorize(enable)`
 - `-Menable-vectorize-pragmas !dir$ ivdep`
 - Tell it that variables or arrays are unique
 - `restrict`
 - Align arrays to cache line boundaries
 - Tell the compiler the arrays are aligned
 - Make loop sizes explicit to the compiler
 - Ensure loops are big enough to vectorise

```
int *loop_size;

void problem_function(float * restrict data1, float * restrict data2, float * restrict data3, int
* restrict index){
    int i,j,n;
    n = *loop_size;
    #pragma ivdep
    for(i=0;i<n;i++){
        j = index[i];
        data1[j] = data2[i] * data3[i];
    }
}
```



Vectorisation example

- This loop doesn't vectorise either:

```
do j = 1,N
  x = xinit
  do i = 1,N
    x = x + vexpr(i,j)
    y(i) = y(i) + x
  end do
end do
```

- Compiler will vectorise inner loop by default
 - Dependency on x between loop iterations


```
do j = 1,N
  x(j) = xinit
end do
do j = 1,N
  do i = 1,N
    x(j) = x(j) + vexpr(i,j)
    y(i) = y(i) + x(j)
  end do
end do
```

Example

```

16.  + 1-----<   do j = 1,N
17.    1             x = xinit
18.  + 1 r4-----<   do i = 1,N
19.    1 r4           x = x + vexpr(i,j)
20.    1 r4           y(i) = y(i) + x
21.    1 r4----->   end do
22.    1----->   end do

```



1.497ms

ftn-6254 ftn: VECTOR File = bufpack.F90, Line = 16

A loop starting at line 16 was **not vectorized** because a recurrence was found on "y" at line 20.

ftn-6005 ftn: SCALAR File = bufpack.F90, Line = 18

A loop starting at line 18 was **unrolled 4 times**.

ftn-6254 ftn: VECTOR File = bufpack.F90, Line = 18

A loop starting at line 18 was not vectorized because a recurrence was found on "x" at line 19.

```

38.  Vf-----<  do j = 1,N
39.  Vf          x(j) = xinit
40.  Vf----->  end do
41.
42.  ir4-----<  do j = 1,N
43.  ir4 if--<    do i = 1,N
44.  ir4 if      x(j) = x(j) + vexpr(i,j)
45.  ir4 if      y(i) = y(i) + x(j)
46.  ir4 if-->    end do
47.  ir4----->  end do

```

x promoted to vector:
trade slightly more memory
for better performance

ftn-6007 ftn: SCALAR File = bufpack.F90, Line = 42

A loop starting at line 42 was **interchanged** with the loop starting at line 43.

ftn-6004 ftn: SCALAR File = bufpack.F90, Line = 43

A loop starting at line 43 was **fused** with the loop starting at line 38.

ftn-6204 ftn: VECTOR File = bufpack.F90, Line = 38

A loop starting at line 38 was **vectorized**.

ftn-6208 ftn: VECTOR File = bufpack.F90, Line = 42

A loop starting at line 42 was **vectorized** as part of the loop starting at line 38.

ftn-6005 ftn: SCALAR File = bufpack.F90, Line = 42

A loop starting at line 42 was **unrolled 4 times**.


1.089ms

-37%



Data alignment

- When vectorising data aligned data is essential for performance



- Unaligned data
 - May require multiple data loads, multiple cache lines, multiple instructions
 - Will generate 3 different versions of a loop: peel, kernel, remainder
- Aligned data
 - Minimum number of data loads/cache lines/instructions
 - Will generate 2 different versions of a loop: kernel and remainder

Aligned data



- Aligned data is best
 - e.g. AVX vector loads/stores operate most effectively on 32-bytes aligned address
 - need to either let the compiler align the data....
 - ..or tell it what the alignment is
- Unit stride through memory is best



Align data



- Align on allocate/create (dynamic)

- `_mm_malloc, _mm_free`
- ```
float *a = _mm_malloc(1024*sizeof(float), 64);
```
- align attribute (at definition, not allocation)
- ```
real, allocatable :: A(1024)
```
- ```
!dir$ attributes align : 64 :: a
```

- Align on definition (static)

```
float a[1024] __attribute__((aligned(64)));
```

```
real :: A(1024)
```

```
!dir$ attributes align : 64 :: a
```

- Common blocks in Fortran

- It's not possible to use directives to align data inside a common block
  - Can align the start of a common block
- ```
!DIR$ ATTRIBUTES ALIGN : 64 :: /common_name/
```
- Up to you to pad elements inside common block

- Derived types

- May need to use `SEQUENCE` keyword and manually pad to get correct alignment



Multi-dimensional alignment



- Need to be careful with multi-dimensional arrays and alignment
 - If you `_mm_malloc` each dimension then it should be fine
 - If you do a single dimension `_mm_malloc` there may be issues:

```
float* a = _mm_malloc(16*15(sizeof(float), 64);  
for(i=0;i<16;i++){  
#pragma clang loop vectorize(enable)  
    for(j=0;j<15;j++){  
        a[i*15+j]++;  
    }  
}
```



Inform on alignment

- For non-static data, as well as aligning data, need to tell compiler it is aligned

- Number of different ways to do this

- Alignment of data inside a loop

- Specify all data in the loop is aligned

```
#pragma vector aligned
!dir$ vector aligned
```

- Alignment of an array

- Specify, for code after the alignment statement, a specific array is aligned

```
__assume_aligned(a, 64);
!dir$ assume_aligned a: 64
```

- May also need to define to properties of loop scalars

```
__assume(n1%16==0);
for(i=0;i<n;i++){
    x[i] = a[i] + a[i-n1] + a[i+n1];
}
!dir$ assume(mod(n1,16).eq.0)
```

- Also can use OpenMP simd clause

- Specify array is aligned for simd loop

```
#pragma omp simd aligned(a:64)
!omp$ simd aligned(a:64)
```

- Different ways of passing data to subroutines can affect performance
- Explicit arrays

```
subroutine vec_add_mult(A, B, C)
real, intent(inout), dimension(1024) :: A
real, intent(in), dimension(1024) :: B, C
```

- Compiler generates subroutine code based on contiguous data
 - Packing/unpacking required to do this is done by the compiler at caller level
 - May be overhead associated with this
- Need to tell the compiler the arrays are aligned (i.e. !dir\$
assume_aligned or !dir\$ vector aligned)
- Same for arrays where array size is passed as an argument to the routine

- Assumed size arrays

```
subroutine vec_add_mult(A, B, C)
real, intent(inout), dimension(:) :: A
real, intent(in), dimension(:) :: B, C
```

- Compiler will generate different versions of the code, with and without contiguous functionality
 - Different versions may show up in the vector reports from the compiler
 - If there are too many different potential versions not all of them will necessarily be generated
 - The fall back version (none unit stride, not vectorised) will be used in this case for inputs that don't match any of the other versions
- Choice which is used made at runtime
- Still need to tell the compiler the arrays are aligned

Fortran data



- Assumed shape arrays

```
subroutine vec_add_mult(A, B, C)
real, intent(inout), dimension(*) :: A
real, intent(in), dimension(*) :: B, C
```

- Compiler generates subroutine code based on contiguous data
 - Packing/unpacking required to do this is done by the compiler at caller level
 - May be overhead associated with this
- Still need to tell the compiler the arrays are aligned



Fortran Indirect addressing

- Indirect addressing code can have some strange affects on vectorisation

```
subroutine vec_add_mult(A, B, C, index)
real, intent(inout), dimension(1024) :: A
real, intent(in), dimension(1024) :: B, C
integer, intent(in), dimension(1024) :: index
integer :: I
```

- Following has flow dependency (needs ivdep directive)

```
do i=1,n
  a(index(i)) = a(index(i)) + b(index(i)) * c(index(i))
end do
```

- Uses gather and scatter operations to pack/unpack indexed locations
- Following creates array temporary for right hand side evaluation

```
a(index(:)) = a(index(:)) + b(index(:)) * c(index(:))
```

- Ends up creating 2 loops

```
temp(:) = a(index(:)) + b(index(:)) * c(index(:))
a(index(:)) = temp(:)
```

- Uses gather/scatter in both loops

OpenMP 4.0 SIMD directives



- Many compilers support their own sets of directives to assist the compiler to vectorise loops.
 - useful but not portable
- OpenMP 4.0 contains a standardised set of directives



Portable SIMD directives



- Use **simd** directive to indicate a loop should be vectorised

#pragma omp simd [*clauses*]

or

!\$omp simd [*clauses*]

- Executes iterations of following loop in SIMD chunks
- Loop is *not* divided across threads
- SIMD chunk is set of iterations executed concurrently by SIMD lanes
- Not a hint! Programmer is asserting independence of iterations.



- Clauses control data environment, how loop is partitioned
- **safelen(length)** limits the number of iterations in a SIMD chunk.
- **linear** lists variables with a linear relationship to the iteration space (induction variables)
- **aligned** specifies byte alignments of a list of variables
- **private**, **lastprivate**, **reduction** and **collapse** have usual meanings.
- Also **declare simd** directive to generate SIMDised versions of functions.
- Can be combined with loop constructs (parallelise and vectorise)
 - **#pragma omp for simd**

Example 1

```
do il=ilg,nm,-1
  jl = il - nm + 1
  if (il-nm.gt.0) then
    d1 = dsqrt(dble((il-nm)*(il+nm)))
    d2 = dble(il)
    ddlf(jl) = (d2*dc*dlf(jl) - d1*dlf(jl-1)) / ds
  else
    ddlf(1) = nm * dc * dlf(1)/ds
  endif
enddo
```

if test prevents vectorisation

Example 1

```
do il=ilg,nm+1,-1
  jl = il - nm + 1
  d1 = dsqrt(dble((il-nm)*(il+nm)))
  d2 = dbble(il)
  ddlf(jl) = (d2*dc*dlf(jl) - d1*dlf(jl-1)) / ds
enddo
ddl(1) = nm * dc * dlf(1)/ds
```

Vectorises!

Example 2

```
do il=2,ilg-nm
  d0 = il+2*nm
  d1 = dble((il-1) * (d0-1))
  d2 = dble(il * d0)
  dbeta = dsqrt(d1/d2)
  d1 = dble(2*(il+nm)-1)
  dalpha = d1/dsqrt(d2)
  dlf(il+1) = dalpha*dlf(il)*dc - dbeta*dlf(il-1)
enddo
```

dependency prevents vectorisation

Example 2

```
do il=2,ilg-nm
  d0 = il+2*nm
  d1 = dble((il-1) * (d0-1))
  d2 = dble(il * d0)
  dbeta(il) = dsqrt(d1/d2)
  d1 = dble(2*(il+nm)-1)
  dalpha(il) = d1/dsqrt(d2)
enddo
```

This loop now vectorizes!

Scalars expanded to arrays.
For large trip counts, loop
blocking would avoid extra
cache misses.

```
do il=2,ilg-nm
  dlf(il+1) = dalpha(il)*dlf(il)*dc - dbeta(il)*dlf(il-1)
enddo
```