

Single Node Optimisation

Optimising with the Compiler



Reusing this material



This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License.

<https://creativecommons.org/licenses/by-nc-sa/4.0/>

This means you are free to copy and redistribute the material and adapt and build on the material under the following terms: You must give appropriate credit, provide a link to the license and indicate if changes were made. If you adapt or build on the material you must distribute your work under the same license as the original.

Note that this presentation contains images owned by others. Please seek their permission before reusing these images.



Motivation

- Why is memory structure important?
 - With current hardware memory access has become the most significant resource impacting program performance.
 - Changing memory structures can have a big impact on code performance.
 - Memory structures are frequently global to the program
 - Different code sections communicate via memory structures.
 - The programming cost of changing a memory structure can be very high.

Programmer's perspective:

- Memory structures are the programmers responsibility
 - At best the compiler can add small amounts of padding in limited circumstances.
 - Compilers can (and hopefully will) try to make best use of the memory structures that you specify (e.g. uni-modular transformations)
- Changing the memory structures you specify may allow the compiler to generate better code.

Types of data structure

- Arrays
- Pointer arrays
- records/structures
- Trees and lists
- Objects

Arrays

- Arrays are large blocks of memory indexed by integer index
- Probably the most common data structure used in HPC codes
- Good for representing regularly discretised versions of dense continuous data

$$f(x, y, z) \rightarrow F[i \mid j \mid k]$$

Arrays

- Multi dimensional arrays use multiple indexes (shorthand)

```
REAL  A(100,100,100)
A (i,j,k) = 7.0
```

```
REAL A(1000000)
A(i+100*j+10000*k) = 7.0
```

```
float  A[100][100][100];
A [i][j][k] = 7.0
```

```
float A[1000000];
A(k+100*j+10000*i) = 7.0
```

- Address calculation requires computation but still relatively cheap.
- Compilers have better chance to optimise where dimension sizes are known at compile time.

Arrays

- Many codes loop over array elements
 - Data access pattern is regular and easy to predict
- Good spatial locality achieved by accessing neighbouring elements on consecutive iterations of the innermost loop.
- Unless loop nest order and array index order match the access pattern may not be optimal for cache re-use.
 - Compiler can potentially address these problems by transforming the loops.
 - But often can do a better job when provided with a more cache-friendly index order.

Bad spatial locality

```
do i=1,n
  do j=1,m
    a(i,j)=a(i,j)+b(i,j)
  end do
end do
```

```
for (j=0; j<M; j++) {
  for (i=0; i<N; i++) {
    a[i][j]+=b[i][j];
  }
}
```

Good spatial locality

```
do j=1,m
  do i=1,m
    a(i,j)=a(i,j)+b(i,j)
  end do
end do
```

```
for (i=0; i<N; i++) {
  for (j=0; j<M; j++) {
    a[i][j]+=b[i][j];
  }
}
```

Dynamic sized arrays (Fortran)

- Not always possible/desirable to fix array sizes at compile time
 - Fortran allows arrays to be dynamically sized based on subroutine arguments.
- Address calculation can still be optimised using CSE.
- Size of slowest moving index is not needed in address computation.
 - Fortran actually allows this dimension to be unspecified in subroutine arguments (assumed size arrays)

Dynamic sized arrays (C)

- C requires array dimensions to be known at compile time.
- However can make slowest dimension variable with pointers and **typedef**

```
typedef float Mat[2][2];  
Mat *data = (Mat *) malloc(n*sizeof(Mat));  
for(i=0; i<n; i++) {  
    for(j=0; j<2; j++) {  
        for(k=0; k<2; k++) {  
            data[i][j][k] = 12.0;  
        }  
    }  
}
```

Pointer arrays

- Alternative to multi-dimensional arrays
 - Pointer to: array of pointers to: array of pointers to: Data

```
float ***data;
data = (float ***) malloc(2*sizeof(float **));
for(i=0;i<2;i++){
    data[i]=(float **) malloc(2*sizeof(float *));
    for(j=0;j<2;j++){
        data[i][j] = (float *) malloc(n*sizeof(float));
        for(k=0;k<n;k++){
            data[i][j][k] = 12.0;
        }
    }
}
```

- Note reverse index order to previous example!

Pointer arrays II

- In C the use-syntax is the same as for arrays
 - `a[i][j][k] = 7.0;`
 - But actually equivalent to
 - `p1 = a[i];`
 - `p2 = p1[j];`
 - `p2[k] = 7.0;`
- Advantage
 - The “columns” are allocated separately and need not be the same length
- Disadvantages
 - Need multiple memory accesses per element access.
 - Need more memory to store all the pointers
 - Less regular access pattern
 - Messy to create/destroy

Records/structures

- Collection of values (of varying types)
 - C structs
 - F90 user defined types
- Good for representing multi-valued data or sparse/scattered data.
- Related variables are stored close together may help cache use.
 - If a code section only uses a subset of the values cache use may suffer.
- Easy to add/re-order members without breaking code as members are referenced by name not position.
 - much harder to remove them.

Structures and the compiler

- Programmer only specifies what a structure contains.
- Compiler chooses layout within the structure.
- In C the compiler usually preserves the order of members but inserts padding between members if needed to meet alignment constraints
 - i.e. Doubles must be aligned on double-word boundaries.
 - Padding reduces cache-line utilisation so order members to reduce padding.
- Similarly in Fortran but can use SEQUENCE keyword to force deterministic layout.

Arrays of structs or structs of arrays?

```
struct Part{  
    double x;  
    double y;  
    double z;  
    int index;  
    double mass;  
}  
Part data[numParts];
```

Array of structs

or

```
struct AllParts{  
    double x[numParts];  
    double y[numParts];  
    double z[numParts];  
    int index[numParts];  
    double mass[numParts];  
}  
AllParts data;
```

Struct of arrays

Array of structs:

- May have good temporal locality if there is lots of computation on each struct
- May have poor spatial locality if computations don't
- Unfavourable for vector loads/stores
- Natural for OO design

Struct of arrays

- May have better spatial locality (use all data on cache line), but worse temporal locality
- More favourable for vector loads/stores
- Less natural for OO design

Arrays of structs of (short) arrays

```
struct FourVecParts{  
    double x[4];  
    double y[4];  
    double z[4];  
    int index[4];  
    double mass[4];  
}  
FourVecParts data[(numParts+3)/4];
```

- Vector friendly without compromising temporal locality too much?
- Not at all natural from a design perspective!

Objects

- Usually implemented much the same as structures
- But objects are opaque
 - Language restricts access to the internal data.
 - Usually need to use special access functions.
- Much easier to change underlying data structure as this is only visible to small fraction of the program
- Access functions introduce additional overhead
 - Function calls
 - Memory copies
- Really only a problem for small low-level objects

Trees/lists

- Structures/Objects can contain pointers to other structures.
 - Can construct trees and lists etc.
- Very flexible and can grow dynamically
 - Same problems as pointer arrays.
 - Additional memory accesses to navigate data
 - Additional storage to store pointers
 - Access pattern is very hard to predict.
- Limited navigation
 - Can only follow access pattern supported by pointer structure
 - e.g. cannot jump to middle of a list without traversing half the nodes.

High level data structures

- Many modern languages have built in-support for high level data structures such as
 - Lists
 - Trees
 - Sets
 - Maps
 - Etc.
- May be available either as built-in data-types or as standard libraries.
 - Have the same intrinsic advantages/disadvantages as home made equivalents but typically better tested and optimised.

What can go wrong

- Poor cache/page use
 - Lack of spatial locality
 - Lack of temporal locality
- Unnecessary memory accesses
 - pointer chasing
 - array temporaries
- Aliasing problems
 - Use of pointers can inhibit code optimisation

Reducing memory accesses

- Memory accesses are often the most important limiting factor for code performance.
 - Many older codes were written when memory access was relatively cheap.
- Things to look for:
 - Unnecessary pointer chasing
 - pointer arrays that could be simple arrays
 - linked lists that could be arrays.
 - Unnecessary temporary arrays.
 - Tables of values that would be cheap to re-calculate.

Utilizing caches

- Want to avoid cache conflicts
 - This happens when too much related data maps to the same cache set.
 - Arrays or array dimensions proportional to (cache-size/set-size) can cause this.
 - Rarely a problem with 8- and 16-way associative caches on XC30
 - Lots of accesses in a loop to arrays with power-of-2 dimensions might still be bad
 - Can pad arrays to avoid this.

Utilizing caches II

- Want to use all of the data in a cache line
 - loading unwanted values is a waste of memory bandwidth.
 - structures are good for this
 - Or loop fastest over the corresponding index of an array.
- Place variables that are used together close together
 - Also have to worry about alignment with cache block boundaries.
- Avoid “gaps” in structures
 - In C structures may contain gaps to ensure the address of each variable is aligned with its size.

Bad Cache Alignment

CrayPAT profiling with export PAT_RT_HWPC=2 (L1 and L2 metrics)

Time%		0.2%
Time		0.000003
Calls		1
PAPI_L1_DCA	455.433M/sec	1367 ops
DC_L2_REFILL_MOESI	49.641M/sec	149 ops
DC_SYS_REFILL_MOESI	0.666M/sec	2 ops
BU_L2_REQ_DC	74.628M/sec	224 req
User time	0.000 secs	7804 cycles
Utilization rate		97.9%
L1 Data cache misses	50.308M/sec	151 misses
LD & ST per D1 miss		9.05 ops/miss
D1 cache hit ratio		89.0%
LD & ST per D2 miss		683.50 ops/miss
D2 cache hit ratio		99.1%
L2 cache hit ratio		98.7%
Memory to D1 refill	0.666M/sec	2 lines
Memory to D1 bandwidth	40.669MB/sec	128 bytes
L2 to Dcache bandwidth	3029.859MB/sec	9536 bytes

cf: 8

Good cache alignment

Time%		0.1%
Time		0.000002
Calls		1
PAPI_L1_DCA	689.986M/sec	1333 ops
DC_L2_REFILL_MOESI	33.645M/sec	65 ops
DC_SYS_REFILL_MOESI		0 ops
BU_L2_REQ_DC	34.163M/sec	66 req
User time	0.000 secs	5023 cycles
Utilization rate		95.1%
L1 Data cache misses	33.645M/sec	65 misses
LD & ST per D1 miss		20.51 ops/miss
D1 cache hit ratio		95.1%
LD & ST per D2 miss		1333.00 ops/miss
D2 cache hit ratio		100.0%
L2 cache hit ratio		100.0%
Memory to D1 refill		0 lines
Memory to D1 bandwidth		0 bytes
L2 to Dcache bandwidth	2053.542MB/sec	4160 bytes

Cache blocking

- A combination of:
 - strip mining (also called loop blocking, loop tiling...)
 - loop interchange
- Designed to increase data reuse:
 - temporal reuse: reuse array elements already referenced
 - spatial reuse: good use of cache lines
- Many ways to block any given loop nest
 - Which loops should be blocked?
 - What block size(s) will work best?



- Analysis can reveal which ways are beneficial
 - How big is your cache?
 - L1 is 512KB on AMD Rome.
 - How many cache lines can it hold?
 - each line typically 64B, so
 - How many cache lines are needed per loop iteration?
 - ...
- But trial-and-error is probably faster
 - or auto-tuning of the code

Loop tiling

```
for (i=0; i<n; i++) {
    for (j=0; j<n; j++) {
        a[i][j]=b[j][i];
    }
}
```



```
for (ii=0; ii<n; ii+=B) {
    for (jj=0; jj<n; jj+=B) {
        for (i=ii; i<ii+B; i++) {
            for (j=jj; j<jj+B; j++) {
                a[i][j]=b[j][i];
            }
        }
    }
}
```



Loop tiling for vectorisation

```
for (i=0;i<n;i++){
    for (j=1;j<n-1;j++){
        a[i][j]=(a[i][j-1] + a[i][j+1])/2.0 ;
    }
}
```

← j loop won't vectorise due to dependencies

```
for (ii=0;ii<n;ii+=B){
    for (j=1;j<n-1;j++){
        for (i=ii;i<ii+B;i++){
            a[j][i]=(a[j-1][i] + a[j+1][i])/2.0 ;
        }
    }
}
```

← i loop will vectorise
but note change of data layout

Further cache optimisations

- If multiple loop nests process a large array
 - First element of array will be out of cache when second loop nest starts
- Improving cache use
 - Consider fusing the loop nests
 - Completely: just have one loop nest
 - Partial: have one outer loop, containing multiple inner loops
 - Beware that too much fusion can result in lots of temporaries and cause the compiler to run out of registers....

Original code	Complete fusion	Partial fusing
<pre>do j = 1, Nj do i = 1, Ni a(i,j)=b(i,j)*2 enddo enddo do j = 1, Nj do i = 1, Ni a(i,j)=a(i,j)+1 enddo enddo</pre>	<pre>do j = 1, Nj do i = 1, Ni a(i,j)=b(i,j)*2 a(i,j)=a(i,j)+1 enddo enddo</pre>	<pre>do j = 1, Nj do i = 1, Ni a(i,j)=b(i,j)*2 enddo do i = 1, Ni a(i,j)=a(i,j)+1 enddo enddo</pre>

Further cache optimisations

- Perhaps cache block before fusing
 - Fuse one or more of the outer blocking loops
- If multiple subprograms process the array
 - Remove one or more outer loops (or all loops) from subprograms
 - Haul loop into parent routine, pass in index values instead
 - Might want to ensure that compiler is inlining this routine
 - This technique is very useful if you want to use OpenMP/OpenACC
- Beware of Fortran
 - array syntax often bad
 - $a(:, :) = b(:, :) * 2$
 - $a(:, :) = a(:, :) + 1$
 - compiler unlikely to fuse any loops

Original code

```
CALL sub1(a,b)
CALL sub2(a)

SUBROUTINE sub1(a)
  do j=1,Nj
    do i=1,Ni
      a(i,j)=b(i,j)*2
    enddo
  enddo
END SUBROUTINE sub1
```

After hauling

```
do j = 1, Nj
  CALL sub1(a,b,j)
  CALL sub2(a,j)
enddo

SUBROUTINE sub1(a,j)
  do i=1,Ni
    a(i,j)=b(i,j)*2
  enddo
END SUBROUTINE sub1
```

Optimising for TLB

- Aim to reuse data on a page
 - i.e. treat similarly to a cache
- Standard-sized pages are 4kB
 - But you can use larger "huge" pages
 - 128kB, 512kB, 2MB,... 64MB
 - Almost always benefit HPC applications
 - regular data accesses
 - huge pages give fewer TLB misses
 - Huge pages can also help communication performance

- To use huge pages (see `man intro_hugepages`)
 - Load chosen `craype-hugepages*` module
 - See `module avail craype-hugepages` for list of available options
 - 2M or 8M are usually most successful on Cray XC30
 - Compile as before
 - Make sure this module is also loaded in PBS jobscript
 - quick cheat: can load a different-sized hugepages module at runtime
 - compile-time module enables hugepages, runtime one determines actual size

Prefetch

- Some processors (including AMD Rome) prefetch automatically
- Regular access patterns are recognised and cache lines fetched in advance.
 - Usually only works for contiguous sequence of cache misses.
- Processor has a set of stream buffers
 - Each holds address of an active stream
 - Loads to the current block causes the next block to be prefetched and the stream address to be updated.
 - Streams are established by series of cache misses to consecutive locations

Using streams

- To utilize stream hardware use linear access patterns where possible
 - Only the order of cache block accesses needs to be linear, not each word access.
- Most loops will require multiple streams
 - If the loop requires more streams than are supported in hardware no prefetching will take place for some of the loads.
 - Consider splitting the loop.
- Prefetching typically cannot cross OS page boundaries
 - huge pages may help

Pointer aliasing

- Pointers are variables containing memory addresses.
 - Pointers are useful but can seriously inhibit code performance.
- Compilers try very hard to reduce memory accesses.
 - Only loading data from memory once.
 - Keep variables in registers and only update memory copy when necessary.
- Pointers could point anywhere, so to be safe compiler will:
 - Reload all values after write through pointer
 - Synchronize all variables with memory before read through pointer

Pointers and Fortran

- F77 had no pointers
- Arguments passed by reference (address)
 - Subroutine arguments are effectively pointers
 - But it is illegal Fortran if two arguments overlap
- F90/F95 has restricted pointers
 - Pointers can only point at variables declared as a “target” or at the target of another pointer
 - Compiler therefore knows more about possible aliasing problems
- Try to avoid F90 pointers for performance critical data structures.

Pointers and C

- In C pointers are unrestricted
 - Can therefore seriously inhibit performance
- Almost impossible to do without pointers
 - malloc requires the use of pointers.
 - Pointers used for call by reference. Alternative is call by value where all data is copied!
- Use the C99 **restrict** keyword where possible
- ...or else use compiler flags
 - CCE: **-h restrict**
 - AMD: **-fstrict-aliasing**
 - GNU: **-fstrict-aliasing -fargument-noalias -fargument-noalias-global**
- Explicit use of scalar temporaries may also reduce the problem